

CLIMB: close the reference-physics loop before allocating policy updates

Tracking difficulty is a gated object

FEASIBILITY $\mathcal{F}(r; M, E)$
robot + scene admissibility

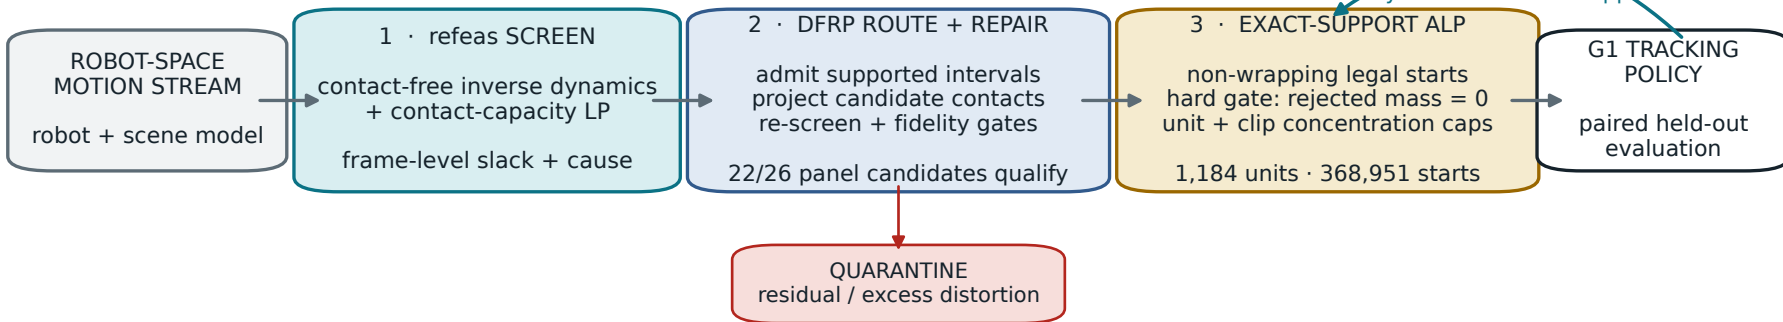
×

SUPPORT $\mathcal{S}(r; B)$
bank-relative representation

×

INTRINSIC $\mathcal{I}(r)$
motion demand

policy outcomes update learning progress
only inside admitted support



Screening changes admission; DFRP recovers qualified data; ALP changes compute allocation. Each interface is measured separately.